

2 **Goodness-of-fit for multi-distribution neutrino** 3 **cross-section measurements with shared events**

4 **F. Martínez López,^{*} and T. Mohayai**

5 *Indiana University, Bloomington, IN 47405, USA*

6 *E-mail:* frmart@iu.edu

7 **ABSTRACT:** When multiple differential cross-section measurements are extracted from a common
8 event sample, the event-sharing structure between distributions imposes exact linear constraints
9 among the bin counts that reduce the effective dimensionality of the measurement below the total bin
10 count. When the statistical covariance matrix accounts for these inter-distribution correlations, the
11 constraints manifest as a rank deficiency. Systematic uncertainties may formally restore invertibility,
12 but the resulting near-singular directions yield unphysically inflated χ^2 values in global goodness-
13 of-fit tests. We present the range-projected χ^2 , a test statistic that restricts the goodness-of-fit
14 test to the subspace carrying independent statistical information, yielding a χ^2 with $N_{\text{bins}} - N_{\text{null}}$
15 degrees of freedom, where N_{null} is the number of independent constraints. We derive N_{null} from
16 the bin-combination structure of the measurement, and show how it decomposes into structural
17 contributions determined a priori from the binning geometry and kinematic contributions that
18 depend on the phase-space occupancy. The method is validated with an analytical toy model and a
19 realistic simulated neutrino–argon cross-section measurement including unfolding and multi-source
20 systematic uncertainties.

21 **KEYWORDS:** Analysis and statistical methods

22 **ARXIV EPRINT:** [1234.56789](https://arxiv.org/abs/1234.56789)

^{*}Corresponding author.

23 Contents

24	1 Introduction	1
25	2 Constraints from shared events	3
26	2.1 Statistical covariance from shared events	3
27	2.2 Combination matrix and null space	4
28	2.3 Determination of the structural null space	5
29	2.4 Effect of systematic uncertainties	6
30	3 Range-projected χ^2 test statistic	8
31	3.1 Test statistic definition	9
32	3.2 Statistical validity	10
33	3.3 Relation to alternative approaches	11
34	4 Toy example	12
35	4.1 Model and measurement setup	12
36	4.2 Null-space structure	14
37	4.3 Goodness-of-fit validation	14
38	4.4 Parameter estimation	16
39	5 Event generator study	17
40	5.1 Model and measurement setup	18
41	5.2 Systematic uncertainties and unfolding	18
42	5.3 Results	18
43	6 Limitations and scope	21
44	7 Summary and conclusions	23
45	A Structural combination matrix algorithm	23

46 1 Introduction

47 Modern neutrino cross-section literature is experiencing an increase of published measurements
48 reporting multiple differential distributions extracted from the same event sample. MicroBooNE
49 has reported simultaneous measurements of cross sections for final states with and without protons
50 [1, 2], multi-differential QE-like cross section measurements in kinematic imbalance variables
51 [3, 4], and triple-differential inclusive results [5]. Other relevant examples include the multiple
52 final-state kinematic projections from T2K [6, 7], and double-differential inclusive results from
53 NOvA [8–10]. In each of these analyses, the same selected events populate bins in multiple

54 reported distributions, leading to strong inter-measurement statistical correlations. Separate but
 55 related efforts have produced simultaneous measurements across different nuclear targets [11–13]
 56 and combined neutrino-antineutrino results [14], where correlations arise from shared systematic
 57 uncertainties rather than shared events.

58 When multiple distributions are extracted from a common event sample, the event-sharing
 59 structure imposes exact linear constraints among the bin counts. If every selected event populates
 60 one bin in each distribution, then the total number of events in all of them is exactly identical. More
 61 generally, the pattern of shared bin occupancy across distributions determines a set of constraints that
 62 reduce the effective dimensionality of the measurement below the total bin count. These constraints
 63 are a property of the measurement itself, independent of any particular statistical treatment.

64 The blockwise unfolding framework of Ref. [15] provides a rigorous method for computing
 65 the full cross-distribution covariance matrices in such measurements. In this approach, all bins
 66 corresponding to the same measured observable in the multi-distribution analysis are organized
 67 in blocks, each corresponding to a single differential cross-section distribution. Each bin block
 68 is unfolded independently using any standard method (e.g. iterative D’Agostini unfolding [16]
 69 or Wiener-SVD [17]), while the full covariance matrix, encoding the correlations between bins
 70 from different distributions, is constructed and propagated through the cross section extraction
 71 procedure. The statistical covariance for bins in different blocks is determined by the number of
 72 events they share. The constraints described above can manifest as a rank deficiency in the statistical
 73 covariance matrix, as certain directions in bin space carry identically zero variance. The number
 74 of these degenerate directions is exactly predictable from the binning geometry.

75 A natural use of these combined measurements is computing global goodness-of-fit metrics,
 76 such as a χ^2 statistic, that test a model against all distributions simultaneously. Compared to
 77 evaluating each distribution independently, a global test that accounts for all inter-block correlations
 78 provides the maximum discriminating power available from the data. Such global metrics are
 79 essential for interaction model tuning efforts [18–20], where measurements of multiple kinematic
 80 distributions provide complementary sensitivity to different aspects of the underlying physics.
 81 However, a global χ^2 that does not account for the constrained directions yields a test with the wrong
 82 effective number of degrees of freedom. When the statistical covariance matrix is constructed under
 83 the Poisson assumption, the constraints appear as exact null directions that prevent direct matrix
 84 inversion. Systematic uncertainties may formally restore invertibility, but the resulting near-singular
 85 directions produce unphysically inflated χ^2 values. In either case, the test is unreliable unless the
 86 constrained directions are properly identified and excluded.

87 We present a method to correctly compute a global goodness-of-fit test statistic for multi-
 88 distribution measurements with shared events. The approach is based on identifying the constrained
 89 directions from the bin structure of the measurement and projecting the statistic onto the subspace
 90 where independent statistical information is available. We refer to the resulting test statistic as the
 91 range-projected χ^2 . We demonstrate the method using a pedagogical toy example and connect it to
 92 the established statistical literature on hypothesis testing with singular covariance matrices.

2 Constraints from shared events

Consider a neutrino cross-section analysis simultaneously measuring B different flux-integrated differential cross sections, each with N_b true bins. If the measurement is extracted from a common event sample then each selected event will populate exactly one bin in each distribution it belongs to. Recent MicroBooNE analyses [21] have adopted the blockwise unfolding approach presented in Ref. [15] to simultaneous cross-section measurements.

This method provides a systematic framework for reporting the correlations between the different extracted differential cross sections. It provides a general prescription to construct statistical covariance matrices that respect the correlations between bins in different distributions or blocks. The full inter-block covariance matrix constructed in reconstructed space is propagated through unfolding using a block-diagonal unfolding matrix, obtained as a direct sum of the individual unfolding matrices computed for each block. The result of the combined measurement is a vector of $N_{\text{bins}} = \sum_{b=1}^B N_b$ cross-section values with an $N_{\text{bins}} \times N_{\text{bins}}$ covariance matrix. The total covariance decomposes as $C_{\text{total}} = C_{\text{stat}} + C_{\text{syst}}$, where the statistical covariance component, computed from the shared-event structure, encodes the constraints described in the introduction.

2.1 Statistical covariance from shared events

The core idea behind the treatment of the statistical correlations in the blockwise unfolding approach is the partition of any arbitrary pair of overlapping kinematic bins X and Y into a set of three non-overlapping bins u , v , and w . The events in the intersection $X \cap Y$ are assigned to bin v , while bin u contains the events in the relative complement of Y in X (i.e. $X \setminus Y$), and bin w corresponds to the relative complement of X in Y (i.e. $Y \setminus X$). Denoting the number of predicted events in each bin as n_i , and assuming these follow independent Poisson distributions, the covariance between the counts in bins X and Y is given by [15]:

$$\begin{aligned} \text{Cov}(n_X, n_Y) &= \text{Cov}(n_u + n_v, n_v + n_w) \\ &= \text{Cov}(n_u, n_v) + \text{Cov}(n_u, n_w) + \text{Cov}(n_v, n_v) + \text{Cov}(n_v, n_w) \\ &= \text{Var}(n_v), \end{aligned} \quad (2.1)$$

as by definition bins u , v , and w have no events in common. This implies that the covariance between any two bins can be estimated from the variance of the events they share.

The variance for a set of Monte Carlo (MC) events can be estimated from the sum of their squared weights. This gives us a general prescription to calculate the elements of the statistical covariance matrix as:

$$(C_{\text{stat}})_{ij} = \sum_{e \in (i \cap j)} w_e^2, \quad (2.2)$$

where w_e is the weight associated to event e belonging to the intersection of bins i and j . In the case of measured data the weights are 1, and the covariance element simplifies to the number of events shared between the bins n_{ij} .

After propagation through the block-diagonal unfolding, the statistical covariance matrix in unfolded space inherits the same rank structure, as the unfolding matrices for each individual block are invertible. This rank structure has a direct interpretation in terms of the number of distinct event-sharing patterns across distributions, which we discuss next.

2.2 Combination matrix and null space

The previous discussion shows that the statistical covariance between any two bins is determined solely by the events they share. We now formalize this observation to derive an exact formula for the rank of C_{stat} , and therefore for the number of directions in bin space that carry zero statistical variance.

For each event e , we can define the indicator vector $\phi_e := (\phi_{e,i})_{i=1}^{N_{\text{bins}}}$ such that $\phi_{e,i} = 1$ if $i \in \text{bins}(e)$ and $\phi_{e,i} = 0$ if $i \notin \text{bins}(e)$, where $\text{bins}(e)$ is the set of bins event e populates across all distributions. The indicator vector acts as an identifier of each event in the combined measurement. For example, in a two-block measurement with 3 bins each, an event populating bin 2 of block 1 and bin 1 of block 2 has indicator vector $\phi_e = (0, 1, 0, 1, 0, 0)$. The statistical covariance matrix can then be re-written in terms of indicator vectors as:

$$C_{\text{stat}} = \sum_e w_e^2 \phi_e \phi_e^T. \quad (2.3)$$

This outer-product decomposition makes the rank structure of C_{stat} clear. The rank of the statistical covariance matrix equals the dimension of $\text{span}\{\phi_e\}$, which depends only on the distinct indicator vectors present, not on their weights. Since all weights are $w_e > 0$ no indicator vector is removed from the sum, and the column space of C_{stat} is exactly $\text{span}\{\phi_e\}$.

We define the bin-combination matrix Ω as the matrix whose rows are the unique indicator vectors among $\{\phi_e\}$. Each row of Ω represent a distinct pattern of bin occupancy across the blocks. By construction the row space of Ω equals $\text{span}\{\phi_e\}$, and therefore:

$$\text{rank}(C_{\text{stat}}) = \text{rank}(\Omega). \quad (2.4)$$

The number of degenerate directions is then:

$$N_{\text{null}} = N_{\text{bins}} - \text{rank}(\Omega). \quad (2.5)$$

Physically, each of the N_{null} directions corresponds to a linear combination of bins along which the measurement has identically zero variance. Any component of the residual vector along such a direction is not perceived by the statistical covariance; it is a direction that the measurement has no statistical power to constrain. Importantly, both Ω and N_{null} are directly computable from the MC event-to-bin mapping alone, before constructing any covariance matrix, confirming that the constraints are a property of the measurement geometry rather than the statistical treatment.

A vector $\mathbf{x}$ is in the null space of C_{stat} if and only if, for every event e :

$$\sum_{i \in \text{bins}(e)} x_i = 0, \quad (2.6)$$

or equivalently $\Omega \mathbf{x} = 0$. Each unique indicator vector acts as a linear constraint on the null-space vectors, requiring that the components of $\mathbf{x}$ sum to zero over the bins populated by the different events. The null space directions can be divided into two physically distinct categories: structural constraints that follow from the block topology alone, and kinematic constraints that depend on which regions of phase space are populated.

2.3 Determination of the structural null space

The combination matrix Ω presented previously is constructed from MC events. One determines which events populate the different bins and collects the unique indicator vectors. However, the indicator vectors that can appear are constrained by two distinct mechanisms. The first is purely structural, the geometry of the binning scheme forbids certain bin combinations. The second is kinematic, correlations between the different variables render some combinations unreachable even though they are geometrically allowed. These constraints are imposed by the underlying physics effects, mainly the differential cross section and the detector acceptance. In the following, we show that the structural constraints can be determined a priori, without generating MC events, from the bin edges used.

The variables used in a measurement are drawn from a set $\{v_1, \dots, v_D\}$. Each block b uses a subset of these variables, $\text{vars}(b) \subseteq \{v_1, \dots, v_D\}$. Restricting the discussion to combinations of single- and double-differential measurements, we have $|\text{vars}(b)| \in \{1, 2\}$ for every b . For each variable v used in a given block b , the binning scheme defines a set of bin edges $E_v^{(b)}$ that partition the range of v into intervals. In the case of a two-dimensional block, the edges in the binning variable may be slice-dependent, but nevertheless these can be collected into a single set of edges per variable across slices. For each variable $v \in \bigcup_b \text{vars}(b)$, we can define the refined edges:

$$\mathcal{E}_v = \bigcup_{b: v \in \text{vars}(b)} E_v^{(b)}, \quad (2.7)$$

with $M_v = |\mathcal{E}_v| - 1$ denoting the number of refined bins for variable v . Since $\mathcal{E}_v$ contains all the bin edges from every block that uses that variable, each refined bin is contained entirely within one original bin of every such block.

The common refinement of all blocks is the Cartesian product of the refined bins across all variables, denoted by $\mathcal{R}$. Each cell $c \in \mathcal{R}$ specifies a refined bin index for every variable in the measurement. For every block b , there exists a unique mapping $\beta_b(c)$ between each cell c and an original bin in said block. The elements of the indicator vector $\phi_c := (\phi_{c,i})_{i=1}^{N_{\text{bins}}}$ associated to cell c can be written as:

$$\phi_{c,i} = \begin{cases} 1 & \text{if } i = \beta_b(c) \text{ for any block } b, \\ 0 & \text{otherwise.} \end{cases} \quad (2.8)$$

If block b does not use variable v then β_b is independent of the refined bin index in v . Therefore, multiple cells may produce identical indicator vectors.

The structural combination matrix Ω_{struct} is defined as the matrix whose rows are the unique indicator vectors from the cells of the common refinement $\{\phi_c : c \in \mathcal{R}\}$. By construction, every indicator vector that appears in the Ω matrix obtained from the MC events is also present in Ω_{struct} , therefore:

$$\text{rank}(\Omega) \leq \text{rank}(\Omega_{\text{struct}}) \leq N_{\text{bins}}. \quad (2.9)$$

This hierarchy motivates a decomposition of the total number of degenerate directions:

$$\begin{aligned} N_{\text{null}} &= N_{\text{bins}} - \text{rank}(\Omega_{\text{struct}}) + \text{rank}(\Omega_{\text{struct}}) - \text{rank}(\Omega) \\ &= N_{\text{null}}^{\text{struct}} + N_{\text{null}}^{\text{kinem}}. \end{aligned} \quad (2.10)$$

The first term, $N_{\text{null}}^{\text{struct}}$, counts null directions that arise purely from the binning geometry and is computable a priori from the bin edges alone. The second term, $N_{\text{null}}^{\text{kinem}}$, accounts for additional null directions introduced by phase-space restrictions that exclude geometrically allowed combinations and can only be determined from Monte Carlo. A description of the algorithm used to compute Ω_{struct} is presented in Appendix A.

The presence of kinematic null directions indicate that certain regions of the combined bin space are empty, which may signal that the binning extends beyond the kinematic reach of the measurement or that it is finer than the available statistics can support. In such cases, merging or removing the offending bins may be preferable to relying on a projection to handle the resulting constraints. The structural null directions, by contrast are an unavoidable consequence of measuring multiple distributions from shared events and cannot be eliminated by rebinning.

The structure of the null space of Ω_{struct} depends on how the bin edges for the different blocks overlap. In the simplest case, when blocks do not share variables, every combination of one bin per block is allowed. The resulting null space is $(B - 1)$ -dimensional, with a single constraint that requires all block totals to be equal. When blocks share one or more variables, the allowed combinations are restricted and the dimension of the null space can be different from $B - 1$. The number of null directions depends on how shared variable bin edges are aligned between blocks. Aligned slice edges decouple the measurement into different sub-problems, adding additional sum rules and increasing $N_{\text{null}}^{\text{struct}}$. On the other hand, misaligned edges couple adjacent distributions, potentially reducing the number of constraints. In the general case, the interplay between shared variables, bin edge alignment, and block topology makes a closed-form expression impractical, but the computation of $\text{rank}(\Omega_{\text{struct}})$ from the bin edge specifications is straightforward.

The structural and kinematic components of the null space respond differently to systematic uncertainties. Whether these cause the rank deficiency of C_{stat} to persist, partially resolve, or fully resolve in the total covariance C_{total} depends on the nature of the systematic treatment, which we discuss next.

2.4 Effect of systematic uncertainties

Previously, we established that the event-sharing structure imposes $N_{\text{null}} = N_{\text{bins}} - \text{rank}(\Omega)$ exact linear constraints on the bin counts, which manifest as a rank deficiency of C_{stat} . Whether the total covariance matrix inherits the rank deficiency depends on which of these null directions receive support from the systematic uncertainties. Now, we examine the mechanisms by which different systematic treatments do or do not populate the null space of C_{stat} .

A systematic variation s produces a modified vector of expected events $\mathbf{n}_s$, whose departure from the central-value prediction can be expressed in terms of a shift vector:

$$\delta_s = \mathbf{n}_s - \mathbf{n}_{\text{CV}}. \quad (2.11)$$

The different contributions to C_{syst} are constructed from the outer product of these shift vectors. A null direction $\mathbf{x}$ of C_{stat} remains a null direction of C_{total} if and only if $\mathbf{x} \cdot \delta_s = 0$ for every systematic variation in the uncertainty budget. The question is therefore whether the shift vectors have nonzero projections onto the structural and kinematic null spaces.

Consider first the simplest case, where a systematic s acts by reweighting the reconstructed event counts directly. Each event e is assigned a modified weight $w_e^{(s)}$ while preserving its bin

231 assignments. The shift vector can then be written as:

$$\delta_s = \sum_e (w_e^{(s)} - w_e) \phi_e. \quad (2.12)$$

232 For any null vector $\mathbf{x}$ which satisfies the null-space condition of Eq. (2.6), the projection of $\mathbf{x}$ onto
233 the shift vector gives:

$$\begin{aligned} \mathbf{x} \cdot \delta_s &= \sum_e (w_e^{(s)} - w_e) (\phi_e \cdot \mathbf{x}) \\ &= \sum_e (w_e^{(s)} - w_e) \sum_{i=1}^{N_{\text{bins}}} \phi_{e,i} x_i \\ &= 0, \end{aligned} \quad (2.13)$$

234 since the inner sum vanishes for every event. Direct reweighting therefore preserves the entire null
235 space, adding no rank to C_{total} . Flux uncertainties, which modify the neutrino energy spectrum
236 through per-event weights, are the classic example. Cross-section model variations and hadronic
237 reinteraction uncertainties, when implemented as per-event reweighting, behave identically.

238 A qualitatively different situation arises when the systematic uncertainties are propagated
239 through the response matrix rather than changing the event counts per bin directly. In this approach,
240 adopted as the standard prescription in MicroBooNE for neutrino interaction uncertainties [15], a
241 modified response matrix Δ'_s , mapping true bins to reconstructed bins under the given variation,
242 is constructed for each systematic uncertainty. This matrix is obtained either by reweighting the
243 events used to build the nominal response matrix or from an independent MC sample with varied
244 simulation parameters. In this case, the predicted reconstructed counts for systematic variation s
245 are:

$$\mathbf{n}_s = \Delta'_s \varphi_{\text{CV}}, \quad (2.14)$$

246 where φ_{CV} is the central-value truth prediction. The shift vector is then:

$$\delta_s = (\Delta'_s - \Delta) \varphi_{\text{CV}}. \quad (2.15)$$

247 Since Δ is block-diagonal, the difference $\Delta'_s - \Delta$ acts independently on each block, projecting
248 the nominal truth vector through block-specific changes in efficiency and migration. The per-event
249 coupling that makes direct reweighting respect all null directions is absent in this prescription. When
250 the weight of an event changes, it affects simultaneously every block it populates; the response-
251 matrix method, in contrast, processes each block independently through its own sub-matrix. The
252 projection of δ_s onto any null vector, whether structural or kinematic, therefore involves differences
253 in how each block responds to the same truth-level prediction, and in general is nonzero. Response-
254 matrix variations thus break both types of null directions, regardless of how Δ'_s is obtained.

255 A third mechanism applies when systematic uncertainties are evaluated by directly comparing
256 reconstructed event counts from independent MC samples simulated with nominal and varied
257 parameters. In this case, the shift vector takes the form:

$$\delta_s = \sum_{e'} w_{e'}^{(s)} \phi_{e'}^{(s)} - \sum_e w_e \phi_e, \quad (2.16)$$

where the first sum runs over the events in the varied-model sample and the second sum over the nominal sample. Since each event in both samples populates exactly one bin per block, their difference projects to zero onto the structural null directions by an analogous argument to that used for the direct reweighting. Kinematic null directions, however, can be lifted if the variation scatters events into reconstructed-bin combinations absent in the nominal sample. In that case, new indicator vectors appear that are not rows of Ω , potentially breaking the corresponding constraints. Detector systematic uncertainties, which usually require independent samples with modified simulation parameters that cannot be captured by event reweighting, are the typical example where this prescription is applied.

In practice, whether the eigenvalues of C_{total} along the formerly null directions are exactly zero or merely smaller than the physical eigenvalues, including these directions in a χ^2 test via standard matrix inversion produces inflated values with the wrong effective number of degrees of freedom. Because the formerly null directions carry no statistical information, C_{total} assigns them a variance set entirely by the systematic model, and the inverse covariance weights them by the reciprocal of that small value. Any residual between data and prediction that projects onto these directions is then amplified, producing a contribution to the χ^2 far in excess of one per degree of freedom. A valid goodness-of-fit test must therefore be restricted to the subspace where the measurement carries genuine sensitivity, which is the role of the projection method developed in the following section.

3 Range-projected χ^2 test statistic

Before constructing the test statistic, we must establish why restricting it to the non-null subspace discards no physical information. Along a null direction $\mathbf{x}$ of C_{stat} the event-sharing constraint $\Omega\mathbf{x} = 0$ holds identically. The same selected events enter the bin combinations on both sides, so the data value along $\mathbf{x}$ is fixed by event conservation rather than measured. Any prediction, forward-folded through the detector response matrix, respects the identical constraint. The residual along $\mathbf{x}$ is therefore zero by construction, up to systematic or numerical excursions that carry no statistical meaning. Therefore, the measurement has no power to constrain a model in this direction. Removing these directions thus removes no genuine sensitivity, only the ill-defined contributions identified in Sec. 2.4. This holds exactly to the extent that the event-sharing structure is exact; when sharing is only approximate the null directions acquire small but nonzero variance, a regime we return to in Sec. 6. Under the exact-sharing assumption, then, it is safe to formulate the global goodness-of-fit test statistic in the subspace orthogonal to the constrained directions, which we now construct.

The constraints are most naturally identified in reconstructed space, where the combination matrix is constructed. The non-null subspace can be obtained decomposing $\Omega^T\Omega$. This matrix is a square $N_{\text{bins}} \times N_{\text{bins}}$ matrix with the same column space and null space as Ω itself, making it a more compact representation of the same information. We can write its eigendecomposition as:

$$\Omega^T\Omega = V S V^T, \quad (3.1)$$

where $S = \text{diag}(s_1, \dots, s_{N_{\text{bins}}})$ contains the eigenvalues in decreasing order and V is the orthogonal matrix of eigenvectors. Of these eigenvalues, $r = \text{rank}(\Omega)$ are nonzero. Partitioning $V = [V_r \mid \tilde{V}_r]$

into the columns associated with the nonzero and zero eigenvalues, the columns of V_r span the subspace of reconstructed bin space that is not constrained by the event-sharing structure. The orthogonal projector onto this subspace is:

$$P_{\text{reco}} = V_r V_r^\top. \quad (3.2)$$

Note that P_{reco} is identical to the projector onto the range of C_{stat} as defined in Eq. (2.3), since C_{stat} and $\Omega^\top \Omega$ share the same column space by construction. The combination matrix formulation makes explicit the fact that the projector is determined by the measurement geometry.

Cross-section measurements are typically reported in unfolded space, and so it is natural to assume this is the space where goodness-of-fit tests should be evaluated. The blockwise unfolding matrix U , which is block-diagonal and invertible, maps reconstructed-space vectors onto unfolded space. The non-null basis vectors propagate as $W_r = U V_r$. Since U is not orthogonal in general, the columns of W_r are not orthonormal, and the projector onto their span takes the general form:

$$P_{\text{true}} = W_r (W_r^\top W_r)^{-1} W_r^\top. \quad (3.3)$$

Since U is invertible, W_r has the same rank as V_r and the number of constrained directions is preserved. However, as shown next, the ordering of projection and unfolding matters as the commutation relation $U P_{\text{reco}} = P_{\text{true}} U$ generically does not hold. For this reason we choose to define the test statistic in reconstructed space, where the constraints originate.

In practice, the eigenvalues of $\Omega^\top \Omega$ that are technically zero take small numerical values. The threshold for identifying these can be set relative to the largest eigenvalue (e.g., $s_i/s_1 < 10^{-10}$). In most cases, a clear spectral gap of several orders of magnitude separates the physical eigenvalues from the numerical noise, making the identification unambiguous.

3.1 Test statistic definition

The goodness-of-fit test is evaluated in reconstructed space, where the residual vector is given by $\delta = \mathbf{d} - \Delta \varphi_{\text{pred}}$, where $\mathbf{d}$ is the background-subtracted data, Δ is the detector response matrix, and φ_{pred} is a truth-level model prediction. The range-projected χ^2 is defined as:

$$\chi_{\text{proj}}^2 = \delta^\top (P_{\text{reco}} C_{\text{total}} P_{\text{reco}})^+ \delta, \quad (3.4)$$

with the superscript $+$ denoting the Moore–Penrose pseudoinverse.

The pseudoinverse can be avoided entirely by reducing the problem to r dimensions. One can verify that the pseudoinverse of $P_{\text{reco}} C_{\text{total}} P_{\text{reco}}$ takes the closed form:

$$(P_{\text{reco}} C_{\text{total}} P_{\text{reco}})^+ = V_r (V_r^\top C_{\text{total}} V_r)^{-1} V_r^\top. \quad (3.5)$$

Substituting into Eq. (3.4) and defining the reduced residual $\tilde{\delta} = V_r^\top \delta$ and the reduced covariance matrix $\tilde{C}_{\text{total}} = V_r^\top C_{\text{total}} V_r$, the test statistic becomes:

$$\chi_{\text{proj}}^2 = \tilde{\delta}^\top \tilde{C}_{\text{total}}^{-1} \tilde{\delta}, \quad (3.6)$$

which involves only a standard inversion of the $r \times r$ matrix $\tilde{C}_{\text{total}}$. Under the null hypothesis, χ_{proj}^2 follows a χ^2 distribution with $r = N_{\text{bins}} - N_{\text{null}}$ degrees of freedom.

Without the projection, the χ^2 test statistic takes the same numerical value whether computed in reconstructed or unfolded space, as the factors of the unfolding matrix cancel identically when applying the additional smearing matrix to the predictions [17]. The introduction of the projector breaks this equivalence when the projection is performed after unfolding. In unfolded-space the residual is defined as $\mathbf{r} = \hat{\boldsymbol{\varphi}} - A_C \boldsymbol{\varphi}_{\text{pred}}$, where $\hat{\boldsymbol{\varphi}} = U\mathbf{d}$ is the unfolded background-subtracted data and $A_C = U\Delta$ is the additional smearing matrix. Using the propagated basis $W_r = U V_r$, one can construct an analogous test statistic in unfolded space as:

$$\chi_{\text{unf}}^2 = \tilde{\mathbf{r}}^T \tilde{V}_{\text{total}}^{-1} \tilde{\mathbf{r}}, \quad (3.7)$$

where $V_{\text{total}} = U C_{\text{total}} U^T$ is the total covariance matrix in unfolded space, and $\tilde{\mathbf{r}} = W_r^T \mathbf{r}$ and $\tilde{V}_{\text{total}} = W_r^T V_{\text{total}} W_r$ are the . While this is a valid test statistic on its own, it is not numerically equal to Eq. (3.6). Expressing it in terms of reconstructed-space quantities yields:

$$\chi_{\text{unf}}^2 = \boldsymbol{\delta}^T U^T U V_r (V_r^T U^T U C_{\text{total}} U^T U V_r)^{-1} V_r^T U^T U \boldsymbol{\delta}, \quad (3.8)$$

which reduces to Eq. (3.6) only when $U^T U = I$, i.e. when the unfolding matrix is orthogonal.

The χ^2 equivalence is restored if the projection is applied before unfolding. Projecting the reconstructed-space residual and covariance first, then propagating through the unfolding matrix, gives:

$$\begin{aligned} (U P_{\text{reco}} \boldsymbol{\delta})^T (U P_{\text{reco}} C_{\text{total}} P_{\text{reco}} U^T)^+ (U P_{\text{reco}} \boldsymbol{\delta}) &= \tilde{\boldsymbol{\delta}}^T W_r^T (W_r \tilde{C}_{\text{total}} W_r^T)^+ W_r \tilde{\boldsymbol{\delta}} \\ &= \tilde{\boldsymbol{\delta}}^T G (G \tilde{C}_{\text{total}} G)^{-1} G \tilde{\boldsymbol{\delta}} \\ &= \chi_{\text{proj}}^2, \end{aligned} \quad (3.9)$$

where the symmetric matrix $G = W_r^T W_r$ acts as the unfolding matrix of the reduced system, and cancels similarly to the unprojected case. In the approach of Eq. (3.8), the full residual, including the null-space components, is first rotated by U into directions that are no longer orthogonal to the range. The unfolded-space projector then picks up this null-space contamination. Projecting first in reconstructed space avoids this by removing the null directions before unfolding can mix them with the physical subspace.

For this reason, Eq. (3.6) is the recommended definition of the range-projected χ^2 .

3.2 Statistical validity

Under the null hypothesis, the residual $\mathbf{r}$ is drawn from a multivariate Gaussian distribution with mean zero and covariance V_{total} . The reduced residual $\tilde{\mathbf{r}} = W_r^T \mathbf{r}$ is a linear transformation of a Gaussian vector, therefore Gaussian itself with mean zero and covariance $\tilde{V}_{\text{total}} = W_r^T V_{\text{total}} W_r$. Provided that $\tilde{V}_{\text{total}}$ is positive definite, guarantees that $\tilde{\mathbf{r}}^T \tilde{V}_{\text{total}}^{-1} \tilde{\mathbf{r}} \sim \chi^2(r)$. The positive definiteness of $\tilde{V}_{\text{total}}$ requires V_{total} to have nonzero variance along every direction in the range of the projector P_{true} , which is generically satisfied. This can be verified by checking that all eigenvalues of $\tilde{V}_{\text{total}}$ are strictly positive.

The validity of the range-projected χ^2 can also be understood from the general theory of quadratic forms of correlated normal variables. Theorem 9.2.2 of Ref. [22] establishes that if $\mathbf{Y} \sim \mathcal{N}(0, \Sigma)$ with $\text{rank}(\Sigma) = k$, then $\mathbf{Y}^T \Sigma^+ \mathbf{Y} \sim \chi^2(k)$ for any choice of generalized inverse Σ^+ .

Assuming the null hypothesis, the projected residual $P_{\text{true}}\mathbf{r}$ is normally distributed with covariance $P_{\text{true}}V_{\text{total}}P_{\text{true}}$, which has rank r by construction. The formulation of Eq. (3.4) with the Moore–Penrose pseudoinverse is a specific instance of this general result. The key distinction from purely numerical approaches to handling singular covariance matrices is that the rank is not determined by an eigenvalue threshold applied to the covariance, but is predicted exactly via the combination matrix Ω from the structure of the measurement.

3.3 Relation to alternative approaches

The range-projected χ^2 is, as a statistical object, an instance of the standard treatment of quadratic forms with singular covariance matrices described in Sec. 3.2; a generalized inverse restricted to the range of the covariance, distributed as a χ^2 with degrees of freedom equal to the rank [22]. Dropping zero-variance directions from a χ^2 is common practice in fits with degenerate covariance matrices. The contribution of this work is not the statistic itself but the recognition that the rank deficiency arises specifically and predictably from the shared-event structure, together with a prescription to compute its dimension a priori from the binning geometry. We discuss here how this differs from related approaches.

The most direct alternative is to regularize the covariance matrix numerically, for instance by performing a singular value decomposition and discarding eigenvalues below a fixed threshold relative to the largest. When the rank deficiency is exact this is unambiguous, but the systematic uncertainties of a realistic measurement lift the null eigenvalues to small but nonzero values, and these lifted directions need not be separated from the genuine physical eigenvalues by a clean gap. A threshold must then be chosen to sit between the lifted null directions and the smallest physical ones, with no guarantee that such a value exists. Setting it too low retains the inflating directions discussed next in Sec. 4.3, while setting it too high discards directions carrying real information. The combination-matrix approach avoids this choice. In this case the rank is fixed by Ω from the measurement geometry, independently of the systematics-contaminated spectrum of C_{total} . A threshold is still used to identify the null space, but it is applied to $\Omega^T\Omega$, where the gap is clean by construction, rather than to C_{total} where systematics have filled it in.

The projection is distinct from, and complementary to, the regularization performed in unfolding. Methods such as Wiener-SVD [17] address the ill-posedness of inverting the response matrix, producing the additional smearing matrix A_C that is applied to truth-level predictions before comparison. They do not address the event-sharing rank deficiency, which is a property of the statistical covariance and survives into unfolded space regardless of the unfolding method. The two operate on different sources of degeneracy and are used together. The event generator study of Sec. 5 unfolds each block with Wiener-SVD and applies the range projection on top.

A more complete alternative is to extract the cross section directly in a binned maximum likelihood fit, in which the truth-level bins are free parameters of interest and the systematic uncertainties enter as nuisance parameters with prior constraints (see, e.g., Ref. [23]). The likelihood is evaluated in reconstructed space, where the inter-distribution correlations are carried by the shared events and the common nuisance parameters of the forward model; it never forms an inverted covariance over the bin counts, so the constrained directions are absorbed automatically and the goodness of fit can be assessed without confronting the rank deficiency at all. This approach is also more robust to the imperfect-sharing scenarios discussed in Sec. 6, where the constraints

are only approximate. Its requirements are correspondingly greater: it needs the full forward model and nuisance parameterization, which are generally available only within the originating group. In this regard, the range-projected χ^2 requires far less information.

For the group producing a blockwise-unfolded measurement it is the natural global goodness-of-fit statistic, computed directly from the covariance and response matrices already used in the extraction, and evaluated in reconstructed space as recommended in Sec. 3.1. It applies equally to an already-published measurement, as performed by generator comparison and tuning frameworks [24]. A standard data release reporting the unfolded data, unfolded covariance, and the block-diagonal unfolding matrix is sufficient to recover the reconstructed-space data and covariance by inversion; the reconstructed-space test then additionally requires the detector response matrix to forward-fold a model prediction. Where the response matrix is not available, the projection can instead be performed in unfolded space using the additional smearing matrix A_C , which is an increasingly standard release product. This unfolded-space test is less exact, as the non-orthogonal unfolding mixes a residual null-space contamination into the physical subspace, but it still removes the bulk of the inflation. The two approaches are complementary: profiling is more general when the complete model is available, while the projection provides a lightweight, a priori test from published results.

4 Toy example

The formalism developed previously is now illustrated with a toy model that captures the essential features of a multi-distribution neutrino cross-section measurement: a two-dimensional underlying physics process, a simple detector model, and a multi-block binning scheme in which events are shared across distributions. The toy is design to be minimal but non-trivial. A three-block structure produces a rank-deficient statistical covariance matrix, and a Feldman–Cousins study demonstrates the consequences of ignoring or correctly treating the resulting null directions.

4.1 Model and measurement setup

The toy model mimics a muon neutrino charged-current quasi-elastic-like interaction, producing a muon characterized by its momentum p_μ and scattering angle θ_μ . The joint truth-level differential cross section is parametrized as:

$$f(p_\mu, \cos \theta_\mu) = p_\mu^k e^{-k p_\mu/p_0} e^{\kappa \cos \theta_\mu} (1 + \alpha p_\mu \cos \theta_\mu), \quad (4.1)$$

where the four parameters (k, p_0, κ, α) control the shape of the momentum spectrum, the level of forward peaking, and the strength of the momentum–angle correlation, respectively. A central-value (CV) model with $(k, p_0, \kappa, \alpha) = (3.0, 0.2, 2.0, 0.3)$ is used to build the response matrix and estimate the covariance, while a “fake” model with $(k, p_0, \kappa, \alpha) = (3.3, 0.2, 1.8, 0.35)$ serves as the true underlying distribution from which pseudo-data are drawn. Truth events are generated via rejection sampling over $p_\mu \in [0.0, 1.2]$ GeV/ c and $\cos \theta_\mu \in [-1.0, 1.0]$.

A simple detector model introduces two different effects. Momentum is smeared multiplicatively, $p_{\text{reco}} = p_{\text{true}}(1 + \delta_p)$ with $\delta_p \sim \mathcal{N}(0, \sigma_p)$ and $\sigma_p = 0.05$, while the angle is smeared additively, $\theta_{\text{reco}} = \theta_{\text{true}} + \delta_\theta$ with $\delta_\theta \sim \mathcal{N}(0, \sigma_\theta)$ and $\sigma_\theta = 1^\circ$. A sigmoid efficiency function




Figure 1. Predicted event distributions in the toy analysis blocks for the CV model (filled histograms) and the fake model (solid lines), together with a pseudo-data realization (points with statistical error bars). The shaded bands indicate the systematic uncertainty on the CV prediction. The top row shows the p_μ and $\cos \theta_\mu$ marginals; the bottom row shows the three $\cos \theta_\mu$ slices of the 2D block, each binned in p_μ .

438 $\varepsilon(p_{\text{reco}}) = [1 + e^{-(p_{\text{reco}} - p_{\text{thr}})/w}]^{-1}$ with threshold $p_{\text{thr}} = 0.10 \text{ GeV}/c$ and width $w = 0.02 \text{ GeV}/c$
 439 models the turn-on of the event acceptance at low momentum.

440 The analysis uses three bin blocks in reconstructed space: a p_μ marginal (5 bins), a $\cos \theta_\mu$
 441 marginal (5 bins), and a two-dimensional distribution sliced in $\cos \theta_\mu$ (3 slices, each with 4 bins
 442 in p_μ), for a total of $N_{\text{bins}} = 22$. The bin edges of the 2D block do not align exactly with those of
 443 the 1D marginals, which is the usual situation in real analyses. A CV sample of $N_{\text{MC}} = 500,000$
 444 truth events is generated, processed through the detector model, and used to construct the response
 445 matrix Δ . Pseudo-data samples are generated by drawing $N_{\text{data}} = 5,000$ events from the fake model.

446 Systematic uncertainties are estimated following the response-matrix approach described in
 447 Sec. 2.4. Interaction-model variations are obtained by reweighting the CV sample to six alternative
 448 model configurations, each producing a varied response matrix Δ_s . The shift vector $\delta_s = (\Delta_s - \Delta) \varphi$
 449 captures the effect on the predicted reconstructed-space distribution. Detector variations are
 450 obtained by generating eight independent samples with perturbed detector parameters (σ_p , σ_θ , p_{thr} ,
 451 and w), each producing a new response matrix from which a shift vector is computed similarly. The
 452 total systematic covariance matrix is $C_{\text{syst}} = N_{\text{var}}^{-1} \sum_i \delta_i \delta_i^T$, scaled to data statistics. The statistical
 453 covariance is computed using Eq. (2.2), assuming unit event weights, which captures the inter-block
 454 correlations from event sharing.

455 Figure 1 shows the predicted event distributions for the CV and fake models in all three blocks




Figure 2. Eigenvalue spectra of the statistical covariance matrix C_{stat} (blue circles) and the total covariance matrix $C_{\text{total}} = C_{\text{stat}} + C_{\text{syst}}$ (orange squares), for the toy analysis, sorted in descending order.

of the toy analysis, together with a pseudo-data sample distribution and the systematic uncertainty bands.

4.2 Null-space structure

The structural combination matrix Ω_{struct} is constructed from the binning configuration of the three blocks, following the prescription from Sec. 2.3. For this bin scheme, Ω_{struct} is rank 17, yielding 5 structural null directions. The combination matrix Ω computed from the MC event list also has rank 17, confirming that all 5 null directions are structural. No additional kinematic constraints arise for this choice of physics model and binning.

The eigenvalue spectrum of C_{stat} confirms this. It shows 5 eigenvalues that are numerically zero, separated from the 17 nonzero eigenvalues by a spectral gap of several orders of magnitude. Adding systematic uncertainties constructed via the response-matrix approach lifts the 5 null eigenvalues, making $C_{\text{total}} = C_{\text{stat}} + C_{\text{syst}}$ formally full rank. However, as shown in Fig. 2, the lifted eigenvalues remain considerably smaller than the bulk of the spectrum. Including them in the test statistic introduces poorly calibrated contributions to the χ^2 , as demonstrated next.

4.3 Goodness-of-fit validation

The calibration of the projected and unprojected χ^2 test statistics is assessed with a toy Monte Carlo exercise at the true hypothesis point. A high-statistics simulation of the fake model is used to estimate the event distribution on a fine reconstructed-space grid, defined by the union of all bin edges in p_μ and $\cos \theta_\mu$ across the three blocks, corresponding to the common refinement described in Sec. 2.3. Each toy experiment draws a Poisson-fluctuated total event count and distributes them across the fine grid via multinomial sampling. These events are then projected onto the three analysis blocks, preserving the inter-block correlations from event sharing. Systematic fluctuations are added




Figure 3. Goodness-of-fit validation for toy experiments generated at the true hypothesis. Left: distributions of the unprojected (red) and range-projected (blue) χ^2 statistics, compared with the expected $\chi^2(22)$ and $\chi^2(17)$ distributions (dashed curves). Right: corresponding p-value distributions, with the false rejection rates at the 5% significance level indicated.

as correlated Gaussian throws from the Cholesky decomposition of C_{sys} . The χ^2 is evaluated in reconstructed space, comparing each toy to the forward-folded prediction at the true hypothesis, for both the unprojected statistic, using the direct inverse of C_{total} , and the range-projected statistic of Eq. (3.4).

Figure 3 shows the results from 10^6 toy experiments. The projected χ^2 statistic closely follows the expected $\chi^2(17)$ distribution, with a mean of 17.23 and a standard deviation of 6.01, in good agreement with the expected values of 17 and $\sqrt{2 \times 17} \approx 5.83$. In contrast, the unprojected statistic is shifted to significantly higher values with a mean of 30.59, and a standard deviation of 8.96, far above the expectation for a $\chi^2(22)$ distribution of 22 and $\sqrt{2 \times 22} \approx 6.63$. The right panel shows the corresponding p-value distributions. The projected statistic yields an approximately uniform distribution, with a false rejection rate of 5.7% at the 5% significance level. The unprojected statistic produces a false rejection rate of 32.4%, more than six times the expected value. An analyzer using the unprojected statistic for goodness-of-fit testing would falsely reject the correct model roughly one third of the time.

Writing the residual vector at the true hypothesis point as $\delta = \delta_{\text{rand}} + \mathbf{b}$, where δ_{rand} has covariance C_{total} and $\mathbf{b} \equiv E[\delta]$ is the deterministic data–prediction residual (estimated as the mean over the toy throws), and diagonalizing $C_{\text{total}} = \sum_j \lambda_j \mathbf{v}_j \mathbf{v}_j^T$, the unprojected statistic decomposes as:

$$\chi^2_{\text{unproj}} = \sum_j \frac{(\delta \cdot \mathbf{v}_j)^2}{\lambda_j}, \quad (4.2)$$

with the expected value of each of these contributions given by:

$$E[(\delta \cdot \mathbf{v}_j)^2 / \lambda_j] = 1 + \frac{(\mathbf{b} \cdot \mathbf{v}_j)^2}{\lambda_j}. \quad (4.3)$$

Each eigendirection contributes as a noncentral $\chi^2(1)$ with noncentrality parameter $\psi_j = (\mathbf{b} \cdot \mathbf{v}_j)^2 / \lambda_j$. On the directions overlapping the null space of C_{stat} the eigenvalue λ_j is the small




Figure 4. Mean contribution to the unprojected χ^2 from each eigendirection of C_{total} , averaged over the toy ensemble, versus the null-space overlap.

498 lifted systematic variance, so any residual projecting onto them is amplified into a contribution far
 499 exceeding one.

500 To identify which directions drive the excess in the unprojected statistic, we quantify how much
 501 each eigenvector of C_{total} overlaps the null space of C_{stat} . Using the projector onto the non-null
 502 subspace P_{reco} from Eq. (3.2), we define the overlap coefficient as:

$$O_j = \mathbf{v}_j^T (I - P_{\text{reco}}) \mathbf{v}_j, \quad (4.4)$$

503 which equals to one for a direction lying entirely in the null space and zero for one in the range.
 504 Figure 4 shows the mean contributions to the χ^2 test statistic across the toy ensemble against O_j . The
 505 17 non-null directions cluster around $O_j \approx 0$ with contributions consistent with unity, confirming
 506 that the statistical fluctuation along them is correctly calibrated. The directions carrying the largest
 507 contributions are all those with substantial null-space overlap, where the small lifted eigenvalue
 508 λ_j amplifies any residual projecting onto them. Summing over all directions, the noncentrality
 509 reproduces the observed excess, $\sum_j \psi_j \approx 8.6$, accounting for the unprojected mean of 30.6.

510 4.4 Parameter estimation

511 The impact of using the projected test statistic on parameter estimation is assessed with a two-
 512 dimensional Feldman–Cousins study. A grid of hypotheses is constructed over (κ, α) with $k = 3.3$
 513 held constant. At each grid point, the truth-level prediction is forward-folded through the nominal
 514 response matrix and compared to the toy pseudo-data. The test statistic $\Delta\chi^2 = \chi^2(\kappa, \alpha) - \chi_{\text{min}}^2$
 515 is computed for both the unprojected and projected cases, and confidence regions are constructed
 516 from the resulting $\Delta\chi^2$ surfaces.

517 Figure 5 shows the 68% and 95% confidence level contours from the Asimov dataset for both
 518 test statistics. The projected contours are correctly centered on the true parameter values, while the
 519 unprojected contours are shifted, with the 68% region excluding the true point. The unprojected




Figure 5. Confidence regions in the (κ, α) plane from the Asimov dataset, showing the 68% and 95% CL contours for the range-projected (blue) and unprojected (red) test statistics. The star marks the true parameter values; the squares indicate the mean best-fit points across 10^6 toys.

best fit values are biased, particularly in α . The mean best-fit value across toys is $\hat{\alpha} = 0.22$, far from the true value of 0.35, while the projected statistic recovers $\hat{\alpha} = 0.37$. The bias in κ is milder ($\hat{\kappa} = 1.84$ for unprojected, $\hat{\kappa} = 1.80$ for projected, compared to the true 1.80). This indicates that α has greater overlap with the null directions.

The frequentist coverage confirms these observations. At the 68.3% confidence level, the projected statistic achieves nominal coverage in both the joint two-dimensional region (68.2%) and the one-dimensional profiled intervals (68.1% for κ , 67.9% for α). The unprojected statistic suffers severe undercoverage: 36.5% in the joint region, 44.6% for κ , and 34.1% for α . The undercoverage is worst for α , consistent with its stronger sensitivity to the null-space contamination. The projected contours are larger than the unprojected ones, reflecting the genuine sensitivity of the measurement once the ill-constrained directions are properly removed.

5 Event generator study

The toy study presented in the previous section validates the usage of the range-projected χ^2 test statistic in a controlled setting, with an analytic cross-section model and known input parameters. Now, we demonstrate that the same pathology arises, and the same fix applies, when working with a realistic measurement setup. In this case, we use a real neutrino flux, neutrino event generators, cross-section systematic uncertainties, and unfolding strategy. The latter is a key addition over the toy study, as it allows us to verify numerically the equivalence between the results in reconstructed and unfolded space, established in Sec. 3.1. In addition, we show that the projected test statistic between generator predictions,

5.1 Model and measurement setup

The central-value model is built from GENIE v3.6.2 [25] using the AR23_20i_00_000 tune on an ^{40}Ar target. The fake-data model, which plays the role of the unknown true cross section, is NuWro v25.11 [26] on the same target. Both generators are run with the SBND BNB ν_μ flux over the energy range 0–10 GeV. Since the flux is shared, flux systematic uncertainties cancel by construction and are not considered.

Events are selected using a topological $\text{CC}1\mu1p0\pi$ definition applied identically to both generators. We require exactly one muon, exactly one proton above 40 MeV kinetic energy, and no charged or neutral mesons above 50 MeV in the final state. This topological selection captures CCQE, CCMEC, and CCRES events where the pion was absorbed in final-state interactions.

The detector model is the same as in the toy study. For each candidate event, we apply a Gaussian momentum and angular smearing ($\sigma_p = 5\%$, $\sigma_\theta = 1^\circ$), and a sigmoid efficiency threshold at $p_{\text{thr}} = 100$ MeV with a width $w = 20$ MeV to the primary muon. Thus, this empty bin is merged with its neighbor, leaving $5 + 5 + (3 + 4 + 4) = 21$ bins across the three blocks.

One million GENIE events are generated and used to construct the response matrix and estimate the systematic uncertainties. One million NuWro events are generated separately; 5,000 of these are drawn as pseudo-data, with the remainder serving as the high-statistics truth prediction.

5.2 Systematic uncertainties and unfolding

Interaction-model systematic uncertainties are estimated using the GENIE Reweight framework [25], which provides per-event weights at $\pm 1\sigma$ for each systematic knob. Fifteen knobs are varied, spanning the dominant sources of uncertainty for the $\text{CC}1\mu1p0\pi$ topology: four z -expansion coefficients for the CCQE axial form factor [27], RPA screening and Pauli blocking corrections [28], the resonance axial and vector masses and overall normalization [29], the CCMEC normalization, two non-resonant single-pion background parameters, and three final-state interaction parameters controlling the pion mean free path, absorption fraction, and charge-exchange fraction [30]. Detector uncertainties are estimated similarly, using $\pm 1\sigma$ variations of the four detector parameters (σ_p , σ_θ , p_{thr} , w). The systematic covariance matrices are constructed via the response-matrix approach described in Sec. 2.4.

Each of the three distribution blocks is unfolded independently using the Wiener-SVD approach [17], with the response matrix and total data covariance computed block-by-block. The per-block unfolding matrices are assembled into the block-diagonal matrix U , and the additional smearing matrix A_C is also constructed from the per-block results. The full covariance is propagated through U as $V_{\text{total}} = U C_{\text{total}} U^T$, preserving the cross-block correlations from event sharing. Truth-level predictions are smeared with A_C before comparison to the unfolded data.

5.3 Results

The structural combination matrix Ω_{struct} , built from the three-block binning scheme, has rank 17, yielding $N_{\text{null}}^{\text{struct}} = 21 - 17 = 4$ structural null directions. The full combination matrix Ω , computed from the GENIE event list, reveals one additional kinematic null direction, bringing the total to $N_{\text{null}} = 5$. This is distinct from the bin merging mentioned in Sec. 5.1. In that case we merged a single bin of zero occupancy, whereas the kinematic null is a linear constraint among the populated




Figure 6. Reconstructed-space event distributions for the realistic fake-data study. The GENIE CV prediction (filled histograms) and systematic uncertainty bands (shaded) are compared to the NuWro prediction (red line) and a 5,000-event NuWro pseudo-data realization (points with statistical error bars). The top row shows the p_μ and $\cos \theta_\mu$ marginals; the bottom row shows the three $\cos \theta_\mu$ slices of the 2D block.

bins. It arises because two cells of the common refinement corresponding to high-momentum backward-going muons are not populated by any selected events, a region where the CCQE-like cross section is strongly suppressed. Their absence removes one independent bin combination from Ω , reducing its rank to 16. The projected test statistic therefore has $r = 16$ degrees of freedom.

Figure 6 shows the reconstructed-space event distributions in all three blocks. The GENIE CV prediction and its systematic uncertainty band are compared to the NuWro fake-data realization and the NuWro prediction. The GENIE systematic band does not fully cover the NuWro prediction in several bins, reflecting genuine model differences that the χ^2 test should be able to identify.

Figure 7 shows the corresponding unfolded distributions, with the GENIE and NuWro truth-level predictions smeared through the additional smearing matrix A_C . Per-block χ^2 values, computed using only the block-diagonal covariance, are annotated in each panel. These are well-behaved for the NuWro comparison, confirming that the unfolding correctly recovers the true distribution at the per-block level. The GENIE comparison shows mild tension in some blocks.

Table 1 summarizes the global χ^2 results. The unprojected statistic yields 121.9/21 for the NuWro prediction. This results in a disproportionate rejection of the correct model, entirely driven by the null-space contamination. The projected statistic gives 15.7/16, consistent with the expected $\chi^2(16)$ distribution and confirming that the method is correctly calibrated. This was further validated with 10^5 pseudo-experiments, which confirmed that the projected χ^2 distribution




Figure 7. Unfolded event distributions for the realistic fake-data study. The GENIE (blue) and NuWro (red) truth-level predictions are smeared through the additional smearing matrix A_C for comparison with the unfolded fake data (points with propagated uncertainties). Per-block χ^2 values are annotated in each panel.

Table 1. Global χ^2 values for the NuWro fake data compared to the GENIE and NuWro truth-level predictions. The projected test statistic has $N_{\text{bins}} - N_{\text{null}} = 16$ degrees of freedom; the unprojected has $N_{\text{bins}} = 21$.

Test statistic	GENIE	NuWro
Unprojected χ^2/ndof	252.4 / 21	121.9 / 21
Projected χ^2/ndof	154.4 / 16	15.7 / 16

follows $\chi^2(16)$ while the unprojected distribution does not follow $\chi^2(21)$, consistent with the toy study results of Sec. 4.3.

The projected statistic retains full discriminating power against the wrong model. The comparison to GENIE yields 154.4/16, a strong rejection indicating that the GENIE uncertainty budget cannot accommodate the NuWro prediction. By contrast, the unprojected statistic gives 252.4/21 for GENIE, also a strong rejection but indistinguishable in practice from its 121.9/21 for NuWro. Without the projection, both the correct and incorrect models are rejected, and the test provides no useful discrimination.

Finally, the projected χ^2 was verified to give identical numerical values whether computed in reconstructed or unfolded space, for both the GENIE and NuWro predictions. This confirms the

equivalence established in Sec. 3.1 and validates the practical implementation of the projection in the presence of unfolding.

6 Limitations and scope

The proposed method rests on the assumption that the event-sharing structure is exact. Every selected event populates exactly one bin in each distribution, so the constraints $\Omega \mathbf{x} = 0$ hold identically and the null space of C_{stat} , built according to Eq. (2.2), coincides with that of Ω . Under this assumption the residual along a null direction is zero by construction and projecting it out discards no statistical information, as argued in Sec. 3.

In real measurements, this assumption is sometimes only approximate. Distributions extracted from a common sample may nonetheless use different phase-space limits, apply distribution-dependent background subtractions, or define their signal components differently, so that the populations entering different blocks are not strictly identical. Consider the simplest constraint that two block totals be equal, encoded by the direction $\mathbf{x}$ with entries +1 on the bins of one block and -1 on those of another. Its statistical variance is given by:

$$\mathbf{x}^T C_{\text{stat}} \mathbf{x} = \sum_e w_e^2 (\phi_e \cdot \mathbf{x})^2 = \sum_{e \in \text{one block only}} w_e^2, \quad (6.1)$$

which vanishes only when every event is shared. When a fraction ϵ of events is unshared, the variance is of order ϵN . In this case, the corresponding eigenvalue is lifted from zero into the body of the spectrum by the statistical model itself, before any systematic uncertainty is added. This is a qualitatively different situation from the inflation of Sec. 4.3. There the variance along some directions was exactly zero in C_{stat} and lifted only by the systematic model to a value unrelated to any statistical power, so dividing a residual by it produced an unphysical contribution. A direction carrying a genuine variance ϵN from real event mismatch is not pathological, the test weights it by its true variance and remains calibrated. The risk of imperfect sharing is therefore not renewed inflation but over-projection, i.e. discarding a direction that now carries a small but real amount of information.

The risk is confined to the a priori structural prescription, and it can be checked for. Three objects encode the null space under exact sharing, and they respond differently when sharing is broken. The structural matrix Ω_{struct} depends only on the binning geometry and is blind to event occupancy entirely. When there is a genuine difference in phase-space acceptance between blocks, this can produce bin combinations absent from Ω_{struct} , breaking the hierarchy of Eq. (2.9) so that $\text{rank}(\Omega)$ exceeds $\text{rank}(\Omega_{\text{struct}})$; the effect is visible in Ω but not in Ω_{struct} . The empirical combination matrix Ω depends on which bin combinations are occupied, but not on how many events carry each, so it is blind to a violation caused only by missing events. In that case we can have $\text{rank}(\Omega) = \text{rank}(\Omega_{\text{struct}})$ even though the affected direction has acquired genuine variance. Only the spectrum of C_{stat} reflects that variance, lifting the corresponding eigenvalue out of the null space.

These two failure modes are illustrated with the toy model of Sec. 4. The first is demonstrated by extending the p_μ acceptance of the two-dimensional block beyond that of the p_μ marginal, so that events with $p_\mu > 1.2 \text{ GeV}/c$ populate the former but not the latter. Of the $O(5 \times 10^5)$ generated




Figure 8. Eigenvalue spectra of the statistical covariance matrix C_{stat} for the toy analysis, sorted in descending order, as events are removed at random from one or two blocks, breaking the exact event-sharing structure. Fractions ϵ_1 and ϵ_2 of events are dropped from the p_μ and $\cos \theta_\mu$ marginal blocks, respectively.

events only seven fall in this region, yet their presence is enough to raise $\text{rank}(\Omega)$ from 17 to 18, above $\text{rank}(\Omega_{\text{struct}}) = 17$, and to make the affected block total fall short of the others by exactly those seven events. The rank signature is thus sensitive to even a single forbidden bin combination.

The second is demonstrated by removing a fraction of events from one block at random, so that the shared-event structure is broken without altering the binning. Figure 8 shows the eigenvalue spectrum of C_{stat} as this fraction is increased independently for blocks 1 (p_μ marginal) and 2 ($\cos \theta_\mu$ marginal). The formerly null eigenvalues, exactly zero under perfect sharing, are lifted into the bulk of the spectrum, filling the spectral gap that otherwise identifies the null space cleanly. The pattern set is unchanged, so $\text{rank}(\Omega) = \text{rank}(\Omega_{\text{struct}})$ throughout; only C_{stat} reflects the violation.

These complementary sensitivities provide a practical set of consistency checks rather than a reason to abandon the geometric construction. Before relying on the structural projection, one should verify that the spectral gap of C_{stat} is clean, that any difference between $\text{rank}(\Omega)$ and $\text{rank}(\Omega_{\text{struct}})$ is understood, and that the measured block totals are consistent within their correlated statistical uncertainties. A deficit, $\text{rank}(\Omega) < \text{rank}(\Omega_{\text{struct}})$, is expected when allowed bin combinations are simply unpopulated; this is the kinematic null case seen in Sec. 5, handled correctly by the empirical Ω . An excess, $\text{rank}(\Omega) > \text{rank}(\Omega_{\text{struct}})$, instead signals that events occupy structurally forbidden combinations, the hallmark of broken sharing. When the diagnostics are consistent with exact sharing, the a priori count of Sec. 2 applies as derived. When they are not, the structural null space over-counts the constrained directions, and projecting onto it would discard information; this signals that the binning or selection definitions should be reconciled across blocks, or that a treatment retaining the affected directions, such as marginalizing over nuisance parameters describing the differences (Sec. 3.3), is more appropriate.

668 7 Summary and conclusions

669 Acknowledgments

670 A Structural combination matrix algorithm

671 We describe the algorithm implemented to construct the structural combination matrix Ω_{struct} from
 672 the bin edge specifications of the multi-block measurement. Each row of Ω_{struct} is a unique indicator
 673 vector encoding one structurally allowed combination of bins, one per block, that a single event
 674 can simultaneously populate. The rank of this matrix determines the number of null directions, as
 675 described in Sec. 2.3.

676 The construction proceeds in three stages. First, the kinematic variables used across all
 677 blocks are classified as shared (appearing in two or more blocks) or private (appearing in exactly
 678 one block). This distinction is key to the efficiency of the algorithm. Shared variables create
 679 correlations between blocks, while private variables contribute bins that are independent of all
 680 other blocks.

681 Second, a common refinement is computed for each shared variable by taking the union of all
 682 bin edges from every block that uses it. The algorithm then enumerates the cells of this shared-
 683 variable refinement. For each cell, the midpoint in each shared variable is used to identify which
 684 bin (or set of bins) in each block is compatible with that cell. When the binning of a given block is
 685 fully determined by the shared variables (e.g., a 1D block whose variable is shared, or a 2D block
 686 whose slice and bin variables are both shared) the cell maps to a single bin. When a block uses
 687 a private variable, the shared cell constrains only part of the binning, and the block contributes
 688 multiple candidate bins corresponding to the different private-variable choices. A 1D block whose
 689 variable is private contributes all of its bins; a 2D block with a shared slice variable and a private
 690 bin variable contributes all bins within the determined slice.

691 Third, the candidate bins from each block are combined. In the simplest case, the full Cartesian
 692 product of all candidate bins is expanded, and each element becomes a row of Ω_{struct} . However,
 693 when many blocks have private variables, this product can become very large. In that case,
 694 a representative subset of indicator vectors is used instead. A base combination is formed by
 695 selecting the first candidate bin from each block, and then the remaining candidates from each block
 696 are substituted one at a time while the others are held at their base values. This produces at most
 697 $1 + \sum_b (n_b - 1)$ rows per shared cell, where n_b is the number of candidate bins for block b . The
 698 representative set spans the same column space as the full Cartesian product, and therefore yields
 699 the same $\text{rank}(\Omega_{\text{struct}})$ and null-space structure.

700 To understand why, note that any indicator vector from the full product can be written as the
 701 sum of indicator vectors from the representative set. Denoting the base combination as $\mathbf{e}_0$ and the
 702 single-block variations as $\mathbf{e}_j^{(k)}$ (block k switched to its j -th candidate), any full-product row that
 703 selects candidate j_k in each block satisfies:

$$\mathbf{e}_{j_1, \dots, j_B} = \mathbf{e}_0 + \sum_k (\mathbf{e}_{j_k}^{(k)} - \mathbf{e}_0), \quad (\text{A.1})$$

704 since the blocks occupy disjoint columns. The full product therefore lies in the span of the
 705 representative rows, and the rank is preserved.

When the measurement involves mutually exclusive event selections, blocks belonging to different selections share no events and are processed independently. The algorithm is applied separately to each selection group, and the resulting matrices are vertically concatenated.

The cost of the algorithm is dominated by the enumeration of shared-variable cells, which scales as $\prod_{v \in \text{shared}} M_v$ where M_v is the number of refined bins for shared variable v . Since only shared variables enter this product, measurements with many blocks but few shared variables, such as a set of 1D distributions in different kinematic variables, are handled efficiently regardless of the total number of bins.

References

- [1] MICROBooNE collaboration, *First Simultaneous Measurement of Differential Muon-Neutrino Charged-Current Cross Sections on Argon for Final States with and Without Protons Using MicroBooNE Data*, *Phys. Rev. Lett.* **133** (2024) 041801 [[2402.19281](#)].
- [2] MICROBooNE collaboration, *Inclusive cross section measurements in final states with and without protons for charged-current ν_μ -Ar scattering in MicroBooNE*, *Phys. Rev. D* **110** (2024) 013006 [[2402.19216](#)].
- [3] MICROBooNE collaboration, *First Double-Differential Measurement of Kinematic Imbalance in Neutrino Interactions with the MicroBooNE Detector*, *Phys. Rev. Lett.* **131** (2023) 101802 [[2301.03706](#)].
- [4] MICROBooNE collaboration, *Multidifferential cross section measurements of $\nu\mu$ -argon quasielasticlike reactions with the MicroBooNE detector*, *Phys. Rev. D* **108** (2023) 053002 [[2301.03700](#)].
- [5] MICROBooNE collaboration, *Measurement of three-dimensional inclusive muon-neutrino charged-current cross sections on argon with the MicroBooNE detector*, *Phys. Lett. B* **870** (2025) 139939 [[2307.06413](#)].
- [6] T2K collaboration, *Characterization of nuclear effects in muon-neutrino scattering on hydrocarbon with a measurement of final-state kinematics and correlations in charged-current pionless interactions at T2K*, *Phys. Rev. D* **98** (2018) 032003 [[1802.05078](#)].
- [7] T2K collaboration, *First T2K measurement of transverse kinematic imbalance in the muon-neutrino charged-current single- π^+ production channel containing at least one proton*, *Phys. Rev. D* **103** (2021) 112009 [[2102.03346](#)].
- [8] NOvA collaboration, *Measurement of the double-differential muon-neutrino charged-current inclusive cross section in the NOvA near detector*, *Phys. Rev. D* **107** (2023) 052011 [[2109.12220](#)].
- [9] NOvA collaboration, *Measurement of $d^2\sigma/d|\vec{q}|dE_{\text{avail}}$ in charged current ν_μ -nucleus interactions at $\langle E_\nu \rangle = 1.86$ GeV using the NOvA Near Detector*, *Phys. Rev. D* **111** (2025) 052009 [[2410.05526](#)].
- [10] NOvA collaboration, “Measurement of the double-differential cross section of muon-neutrino charged-current interactions with low hadronic energy in the NOvA Near Detector.” 10, 2024.
- [11] T2K collaboration, *Simultaneous measurement of the muon neutrino charged-current cross section on oxygen and carbon without pions in the final state at T2K*, *Phys. Rev. D* **101** (2020) 112004 [[2004.05434](#)].

- [12] MINERvA collaboration, *Simultaneous Measurement of ν_μ Quasielasticlike Cross Sections on CH, C, H₂O, Fe, and Pb as a Function of Muon Kinematics at MINERvA*, *Phys. Rev. Lett.* **130** (2023) 161801 [2301.02272].
- [13] MINERvA collaboration, *Simultaneous Measurement of Muon Neutrino ν_μ Charged-Current Single π^+ Production in CH, C, H₂O, Fe, and Pb Targets in MINERvA*, *Phys. Rev. Lett.* **131** (2023) 011801 [2209.07852].
- [14] T2K collaboration, *First combined measurement of the muon neutrino and antineutrino charged-current cross section without pions in the final state at T2K*, *Phys. Rev. D* **101** (2020) 112001 [2002.09323].
- [15] S. Gardiner, “Mathematical methods for neutrino cross-section extraction.” 2024.
- [16] G. D’Agostini, *A Multidimensional unfolding method based on Bayes’ theorem*, *Nucl. Instrum. Meth. A* **362** (1995) 487.
- [17] W. Tang, X. Li, X. Qian, H. Wei and C. Zhang, *Data Unfolding with Wiener-SVD Method*, *JINST* **12** (2017) P10002 [1705.03568].
- [18] C. Wilkinson et al., *Testing charged current quasi-elastic and multinucleon interaction models in the NEUT neutrino interaction generator with published datasets from the MiniBooNE and MINERvA experiments*, *Phys. Rev. D* **93** (2016) 072010 [1601.05592].
- [19] MINERvA collaboration, *Tuning the GENIE Pion Production Model with MINERvA Data*, *Phys. Rev. D* **100** (2019) 072005 [1903.01558].
- [20] GENIE collaboration, *Neutrino-nucleus $CC0\pi$ cross-section tuning in GENIE v3*, *Phys. Rev. D* **106** (2022) 112001 [2206.11050].
- [21] MICROBooNE collaboration, *Measurement of single- and double-differential cross sections for mesonless charged-current muon neutrino interactions on argon with final-state protons using the MicroBooNE detector*, *Phys. Rev. D* **112** (2025) 112004 [2403.19574].
- [22] C. Rao and S. Mitra, *Generalized Inverse of Matrices and Its Applications*, Probability and Statistics Series, Wiley (1971).
- [23] T2K collaboration, *First measurement of muon neutrino charged-current interactions on hydrocarbon without pions in the final state using multiple detectors with correlated energy spectra at T2K*, *Phys. Rev. D* **108** (2023) 112009 [2303.14228].
- [24] P. Stowell et al., *NUISANCE: a neutrino cross-section generator tuning and comparison framework*, *JINST* **12** (2017) P01016 [1612.07393].
- [25] C. Andreopoulos et al., *The GENIE Neutrino Monte Carlo Generator*, *Nucl. Instrum. Meth. A* **614** (2010) 87 [0905.2517].
- [26] T. Golan, J.T. Sobczyk and J. Zmuda, *NuWro: the Wroclaw Monte Carlo Generator of Neutrino Interactions*, *Nucl. Phys. B Proc. Suppl.* **229-232** (2012) 499.
- [27] A.S. Meyer, M. Betancourt, R. Gran and R.J. Hill, *Deuterium target data for precision neutrino-nucleus cross sections*, *Phys. Rev. D* **93** (2016) 113015 [1603.03048].
- [28] J. Nieves, I. Ruiz Simo and M.J. Vicente Vacas, *Inclusive Charged-Current Neutrino-Nucleus Reactions*, *Phys. Rev. C* **83** (2011) 045501 [1102.2777].
- [29] D. Rein and L.M. Sehgal, *Neutrino Excitation of Baryon Resonances and Single Pion Production*, *Annals Phys.* **133** (1981) 79.

786 [30] S. Dytman, Y. Hayato, R. Raboanary, J.T. Sobczyk, J. Tena Vidal and N. Vololoniaina, *Comparison of*
787 *validation methods of simulations for final state interactions in hadron production experiments*, *Phys.*
788 *Rev. D* **104** (2021) 053006 [[2103.07535](#)].